



The router gives no specialisation pressure, but the resulting redundancy is NOT expert-level $\Rightarrow$ look inside the expert (Qwen3-30B-A3B, E=128, K=8)

(a) Experts are fed the same kind of token



(b) Their weights do share a subspace



(c) ...yet no expert can stand in for another

